

JAHANVI JESWANI

📍 Seattle, WA 📞 +1 (206) 370-9806 ✉️ jjeswani@uw.edu [LinkedIn](#) [GitHub](#)

Education

University Of Washington

Sep 2025 – Expected Mar 2027

Master of Science (M.S), Data Science | **GPA: 3.92/4.0**

Seattle, WA

- Relevant Courses: Machine Learning, Natural Language Processing, Applied Statistics and Experimental Design, Probability and Statistics, Scalable Systems and Algorithms, Explainable AI

Shri G. S. Institute of Technology and Science

Aug 2018 – May 2022

Bachelor of Technology (B.S), Electronics and Instrumentation | **GPA: 3.78/4.0**

Indore, India

Technical Skills

Languages: Python, SQL, PySpark, R, C++, JavaScript, TypeScript, GIT, Bash

ML Frameworks: PyTorch, TensorFlow, Scikit-learn, XGBoost, HuggingFace, Pandas, NumPy, Matplotlib

Gen AI & LLMs: RAG, LangChain, LangGraph, Embeddings, Multi Agent Orchestration, LoRA, Claude APIs

MLOps & Cloud: MLflow, Databricks, Docker, Git, CI/CD, AWS (SageMaker, S3), Azure, GCP, Streamlit

Experience

Coupage

Jun 2026 – Sep 2026

SDE-AI Intern, Coupage Intelligent Cloud (CIC)

Seattle, WA

- Built an **LLM inference benchmarking framework (GuideLLM)** measuring **throughput**, **TTFT**, **latency**, **cost**, and **quality** across concurrency, long-context, and traffic profiles, feeding a live **Streamlit** dashboard
- Benchmarked **50+ model-stack configurations** for a production API, informing a stack migration that cut **latency 15%** and **cost 4%** while raising concurrency capacity; extended to **GPU burn-in** validation via **MLPerf Training**

HP Inc.

Apr 2023 - Aug 2025

Data Scientist

Bengaluru, India

- Built a **natural language to dashboard** system using **Llama 3** with **schema-aware few-shot prompting**, translating stakeholder queries into **Plotly** visualizations via **Streamlit**, enabling self-serve data exploration
- Built an **XGBoost** propensity model for **HP iPaper** targeting with **A/B testing** and quarterly **feedback loop** retraining, acquiring **1M+ customers** and driving **\$55M** in revenue
- Deployed a revenue forecasting model (**R2 = 0.80**) with quarterly retraining via **Databricks**, **MLflow**, and **Azure CI/CD pipelines**, with a **SHAP**-powered Streamlit simulator for sales partners to model compensation scenarios
- Built scalable **CLV and churn inference pipelines** using **Survival Analysis (Cox PH)** and **CatBoost** on billions of telemetry events, engineering features and optimizing **PySpark** pipelines to reduce runtime by **40%**

Groww

Jun 2022 – Mar 2023

Machine Learning Engineer-1

Bengaluru, India

- Reduced a **3-month KYC backlog** by building a **YOLO & OCR** pipeline to locate ID proofs in unstructured PDFs and cross-validate user-input information against extracted fields, flagging **40%** low-confidence cases for fraud review
- Increased Fixed Deposit adoption by **6–8% QoQ** by building an **XGBoost** propensity model on millions of clickstream interactions, validated through **A/B testing** to identify high-risk portfolio holders as unexpected high-conversion targets

Groww

Jan 2022 – May 2022

Machine Learning Intern

Bengaluru, India

- Clustered **BERT embeddings** across app reviews and email surveys, focusing on low ratings to surface **8 recurring pain points** in the Mutual Funds product, informing UI/UX fixes that reduced support complaints by **15%**

Projects

MultiEval | 2nd Place, LA Hacks 2026 | *Python, TypeScript, Openhands, inspect-ai, Claude, Gemini*

- Built a **multi-agent evaluation platform** on inspect-ai with **6 harnesses** (incl. **EV-gated delegation**, **Cross-Sibling-Discovery**) and a custom **Gantt + causal-graph** viewer with **A/B testing** (Wilson CI, Cohen's d)
- Designed **Structured Parent-Child Task Delegation** harness, lifting pass rate **0.79 → 0.87 (+8pp)** over baseline

LectureAI | *Python, FastAPI, Gemini, Next.js, TypeScript, Tailwind CSS*

- Transformed YouTube lectures into study environments by building a **full-stack 6-agent system** with **FastAPI** and **Next.js**, deployed on Railway and Vercel with **SSE streaming** and **YouTube API v3 fallback** for reliability
- Engineered a **Curriculum mapper** via **Gemini embeddings** to surface coverage gaps across **10+ lecture series**

LLM Alignment & Cross-Cultural Reliability Analysis | *LLM Evaluation, OpenAI API, HuggingFace*

- Quantified Value-Action Gap in **GPT-4o-mini & Gemma-2-9B** across **14.8k social scenarios** and **12 cultures**, improving alignment via **LoRA fine-tuning (SFT + DPO)** with a dataset of aligned vs misaligned decisions

Publications

- Low Valence Low Arousal Stimuli for EEG-Based Biometric Authentication **J. Jeswani**, et al. | **MIE 2023**, Stud Health Technol Inform, IOS Press | PMID: 37203658 [View Publication]