

JAHANVI JESWANI

📍 Seattle, WA 📞 +1 (206) 370-9806 ✉️ jjeswani@uw.edu [LinkedIn](#) [GitHub](#)

Education

University Of Washington

Sep 2025 – Expected Mar 2027

Master of Science in Data Science; GPA: 3.90

Seattle, WA

- **Coursework:** Machine Learning, NLP, Deep Learning, Applied Statistics, Scalable Systems, Explainable AI

Shri G. S. Institute of Technology and Science

Aug 2018 – May 2022

Bachelor of Technology in Electronics and Instrumentation; GPA: 3.78

Indore, India

- **Relevant Courses:** Data Structures, Algorithms, Database Management Systems, OOPS, Statistics, Linear Algebra

Technical Skills

Languages: Python, SQL, PySpark, R, C++

ML Frameworks: PyTorch, TensorFlow, Scikit-learn, XGBoost, HuggingFace, Pandas, NumPy, Keras

Gen AI & LLMs: BERT, GPT-4o, Llama 3, LangChain, RAG, Embeddings, Vector Search, Tool Calling

MLOps & Cloud: MLflow, Databricks, Docker, Git, CI/CD, AWS (SageMaker, S3), Azure, GCP

Experience

HP Inc.

Apr 2023 - Aug 2025

Data Scientist

Bengaluru, India

- Built a **natural language to dashboard** system using **Llama 3** with **schema-aware few-shot prompting**, translating stakeholder queries into **Plotly** visualizations via **Streamlit**, enabling self-serve data exploration
- Built scalable **CLV and churn inference pipelines** using **Survival Analysis (Cox PH)** and **CatBoost** on billions of telemetry events, engineering features and optimizing **PySpark** workflows to reduce runtime by **40%**
- Deployed a **revenue forecasting model (R2 = 0.80)** with **quarterly retraining** via **Databricks**, **MLflow**, and a **SHAP-powered Streamlit** simulator for sales partners to model compensation scenarios
- Built an **XGBoost** propensity model for **HP iPaper** targeting with **A/B testing** and quarterly **feedback loop** retraining, driving **1M+ subscriptions in 1 year** and reducing marketing costs by **50%**

Groww

Jun 2022 – Mar 2023

Machine Learning Engineer-1

Bengaluru, India

- Increased Fixed Deposit adoption by **6–8% QoQ** by building an **XGBoost** propensity model on millions of clickstream interactions, validated through **A/B testing** to identify high-risk portfolio holders as unexpected high-conversion targets
- Clustered **BERT embeddings** across app reviews, ratings, and surveys for the **Mutual Funds** product to surface recurring user pain points, directly informing **UI/UX fixes** that reduced related support complaints

Groww

Jan 2022 – May 2022

Machine Learning Intern

Bengaluru, India

- Built a **YOLO + OCR** pipeline to detect and classify ID documents from unstructured KYC PDFs at **78% precision**, reducing a **3-month backlog** and flagging **40% low-confidence cases** for fraud review and cutting verification costs

Projects

Agentic Customer Service System | Python, LangChain, GPT-4o, PostgreSQL

- Built a production-style **AI agent** using **GPT-4o** with **tool calling** and **memory**, automating customer support workflows (returns, order validation, inventory checks), reducing manual intervention through end-to-end task execution
- Designed for rapid iteration with user feedback loops, with multiple versions to improve usability and task success rates

ML-Powered Accessibility Extension | 1st Place Datathon | XGBoost, SHAP, GPT-4o, JavaScript, FastAPI

- Built a real-time **Chrome extension** with an **end-to-end ML pipeline: XGBoost (0.85 F1)** for violation severity classification, **SHAP** explanations piped into **GPT-4o** to auto-generate targeted fix recommendations via **FastAPI**

Document Q&A via RAG | LlamaIndex, LangChain, OpenAI, Streamlit

- Built an internal **knowledge assistant** using **RAG** (LlamaIndex + LangChain) over **1K+ documents**, enabling teams to retrieve operational knowledge and reducing time spent searching across fragmented sources

LLM Alignment & Cross-Cultural Reliability Analysis | LLM Evaluation, OpenAI API, HuggingFace

- Quantified Value-Action Gap in **GPT-4o-mini & Gemma-2-9B**, identifying a **60% alignment rate** across **14.8k social scenarios** and **12 cultures** to expose systemic inconsistency in **LLM moral priors**

Publications and Awards

- Low Valence Low Arousal Stimuli for EEG-Based Biometric Authentication **J. Jeswani**, et al. | Caring is Sharing: Exploiting the Value in Data for Health, IOS Press (2023) [View Publication]
- **National Finalist**, Philips Code to Care Challenge Season-2: Ranked in the top 8 teams all over India
- **1st Place** – Databricks x UW Hackathon | **1st Place** – UW Datathon 2026 (Technology Access Track)